ASSIGNMENT-3

Applied Data Science

ASSIGNMENT DATE	08 october 2022
STUDENT NAME	Ms.Mythili
STUDENT REG. NO	721719104053
MAXIMUM MARKS	2 Marks

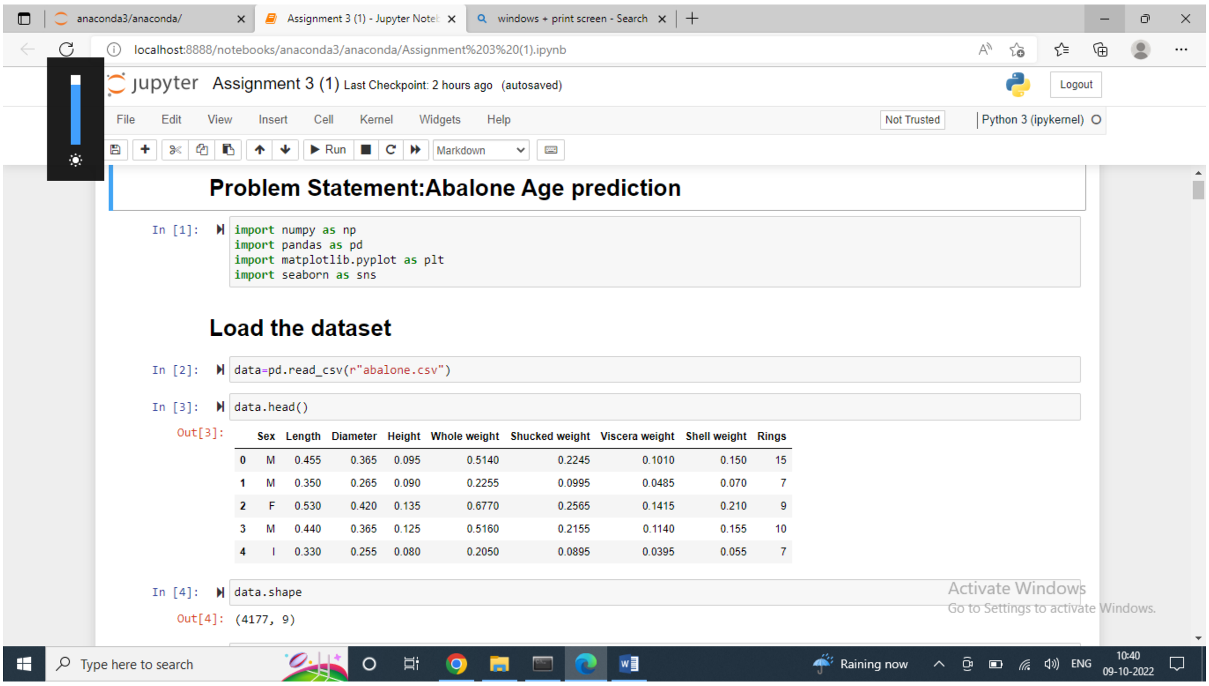

Problem Statement:Abalone Age prediction

In [1]:
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Load the dataset

In [2]:
```
data=pd.read_csv(r"abalone.csv")
```

In [3]:
```
data.head()
```

Out[3]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

In [4]:
```
data.shape
```

Out[4]: (4177, 9)

```
In [5]: data.size
Out[5]: 37593
```

```
In [6]: data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  ------          --------------  -----
 0   Sex             4177 non-null   object
 1   Length          4177 non-null   float64
 2   Diameter        4177 non-null   float64
 3   Height          4177 non-null   float64
 4   Whole weight    4177 non-null   float64
 5   Shucked weight  4177 non-null   float64
 6   Viscera weight  4177 non-null   float64
 7   Shell weight    4177 non-null   float64
 8   Rings           4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

Data Analysis & Visualization

Univariate Analysis

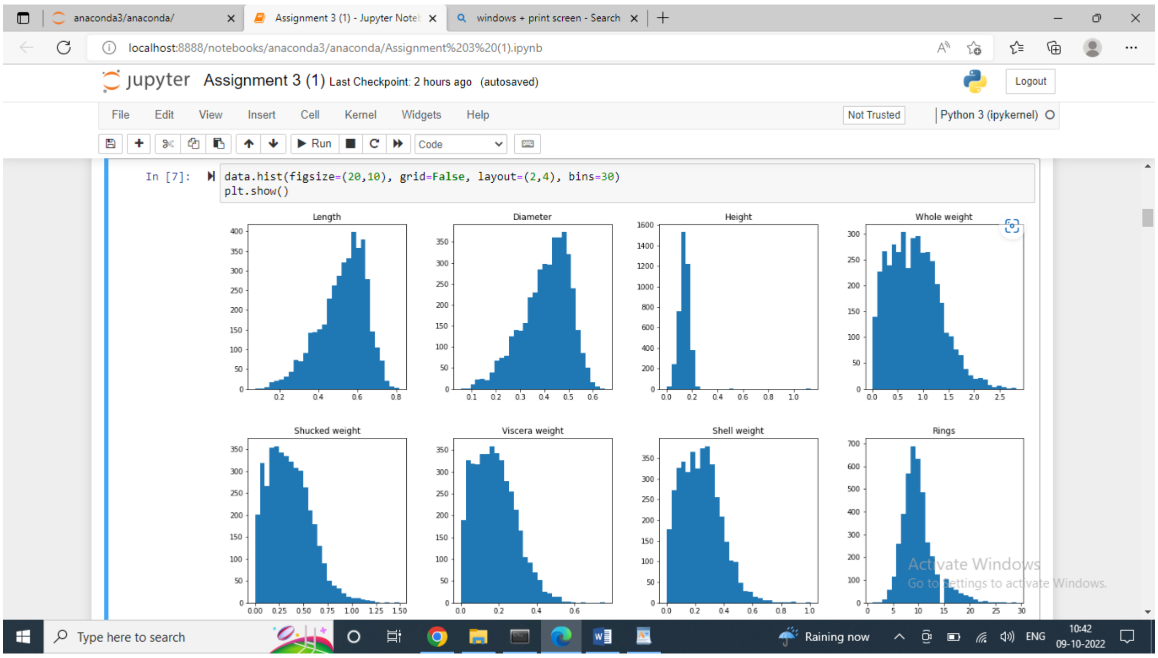

anaconda3/anaconda/
Assignment 3 (1) - Jupyter Note
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
In [7]:
data.hist(figsize=(20,10), grid=False, layout=(2,4), bins=30)
plt.show()
Length
Diameter
Height
Whole weight
Shucked weight
Viscera weight
Shell weight
Rings
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:42
09-10-2022

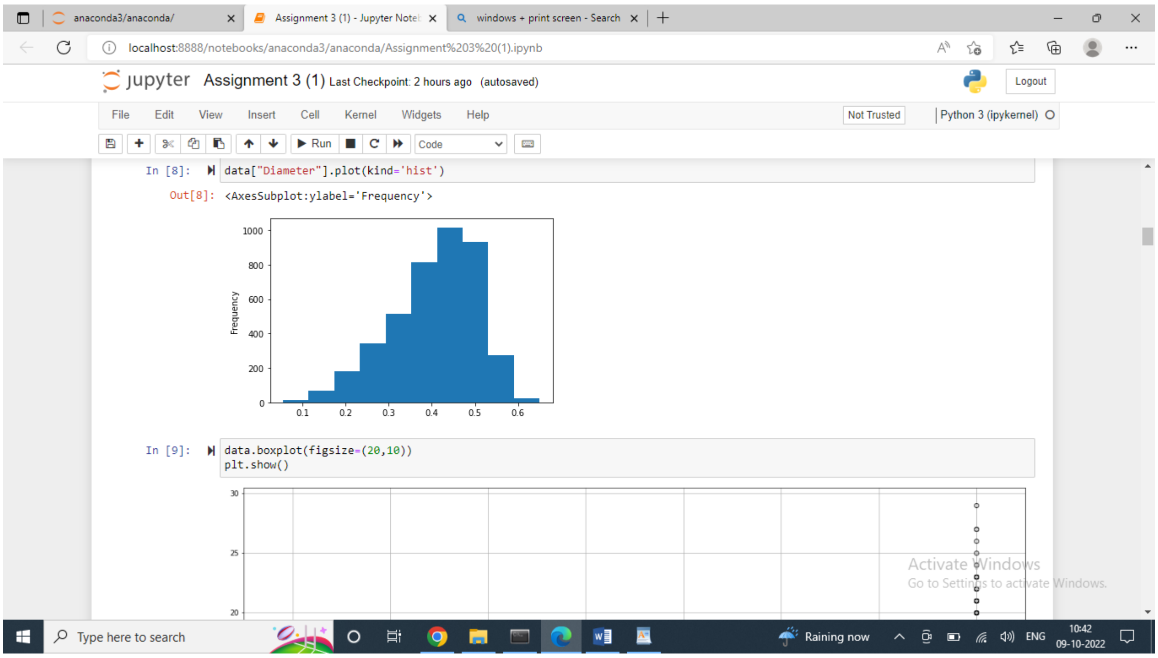
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Notel
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
In [8]: data["Diameter"].plot(kind='hist')
Out[8]: <AxesSubplot:ylabel='Frequency'>
Frequency
1000
800
600
400
200
0
0.1 0.2 0.3 0.4 0.5 0.6
In [9]: data.boxplot(figsize=(20,10))
plt.show()
30
25
20
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:42
09-10-2022

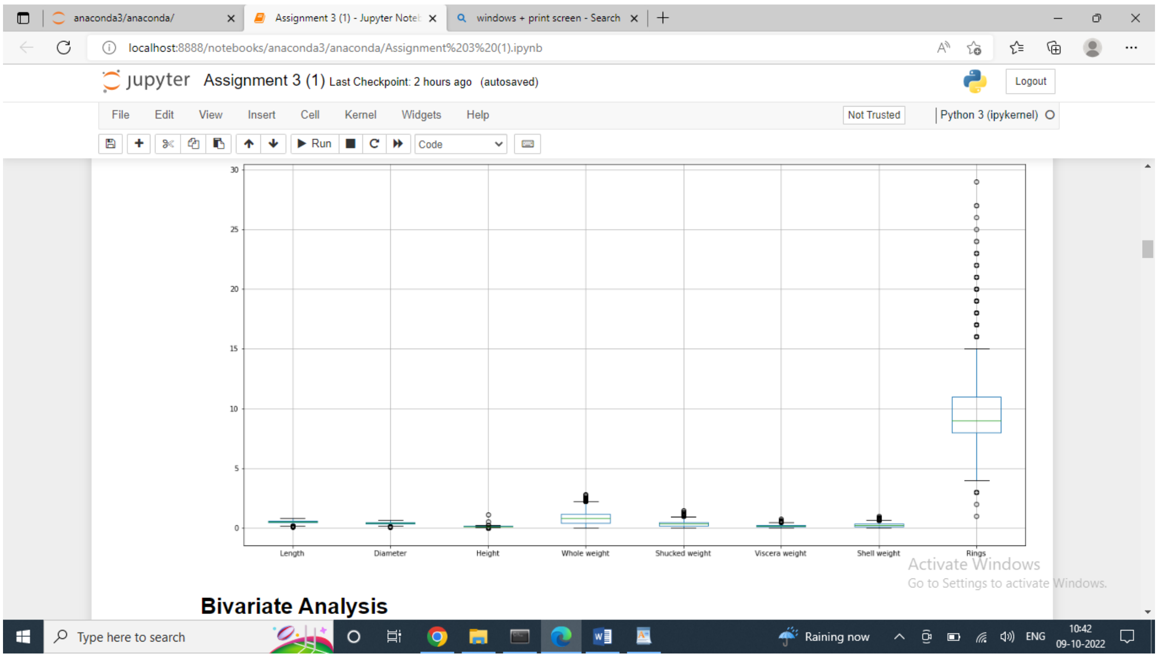
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Noteb
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
30
25
20
15
10
5
0
Length
Diameter
Height
Whole weight
Shucked weight
Viscera weight
Shell weight
Rings
Activate Windows
Go to Settings to activate Windows.
Bivariate Analysis
Type here to search
Raining now
ENG
10:42
09-10-2022

Bivariate Analysis

In [10]:
```
data.head()
```

Out[10]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

In [11]:
```
fig, axes = plt.subplots(4,2, figsize=(15,15))
axes = axes.flatten()
for i in range(1,len(data.columns)-1):
    sns.scatterplot(x=data.iloc[:,i], y=data['Rings'], ax=axes[i])
plt.show()
```

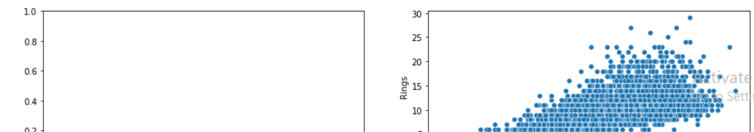

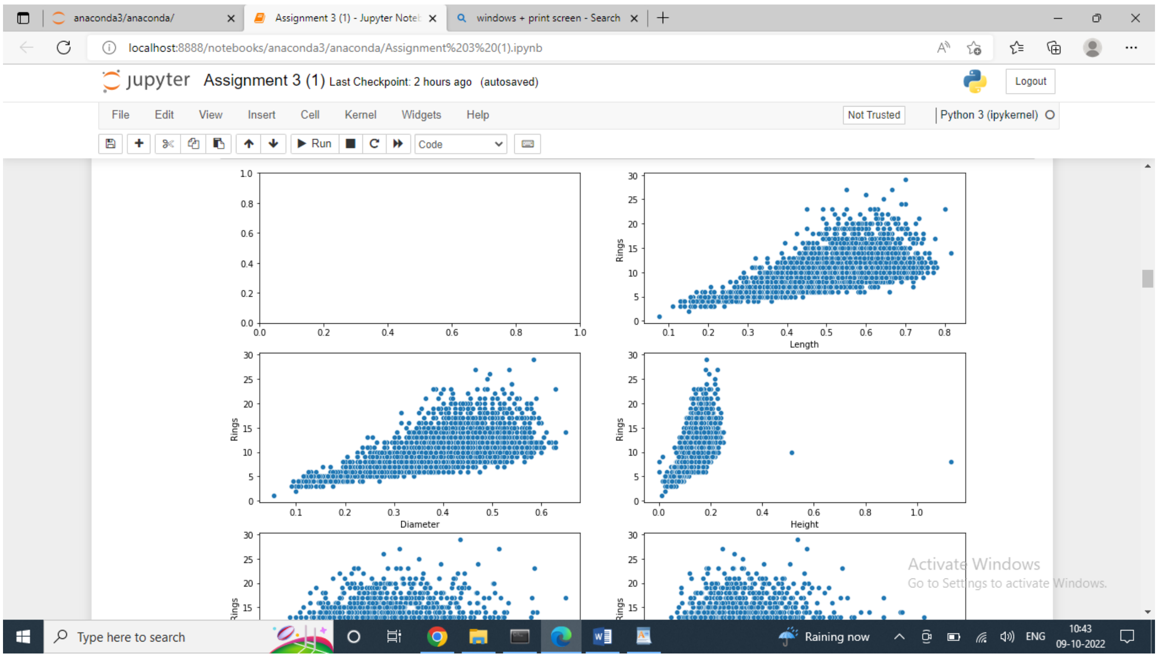
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
Rings
Length
Diameter
Height
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:43
09-10-2022

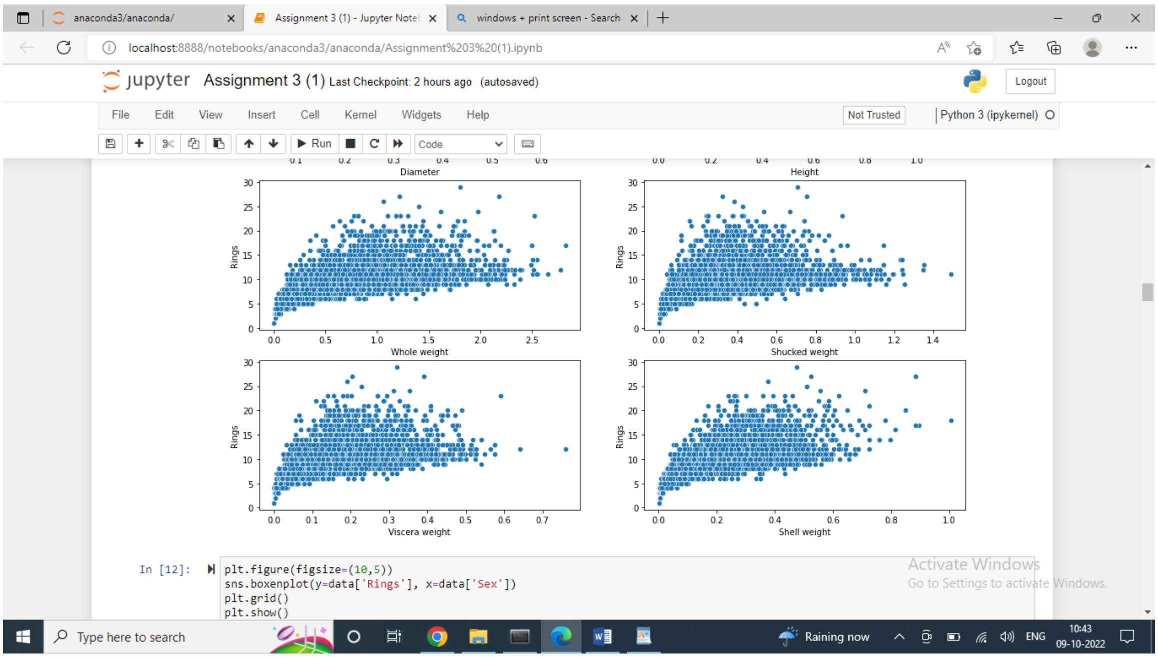
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Noteb
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
Diameter
Height
Rings
Whole weight
Shucked weight
Viscera weight
Shell weight
In [12]:
plt.figure(figsize=(10,5))
sns.boxenplot(y=data['Rings'], x=data['Sex'])
plt.grid()
plt.show()
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:43
09-10-2022

jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)

In [12]:

```
plt.figure(figsize=(10,5))
sns.boxenplot(y=data['Rings'], x=data['Sex'])
plt.grid()
plt.show()
data.groupby('Sex')['Rings'].describe()
```

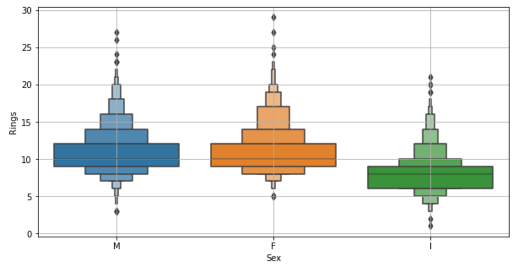

Out[12]:

Sex	count	mean	std	min	25%	50%	75%	max
F	1307.0	11.129304	3.104256	5.0	9.0	10.0	12.0	29.0
I	1342.0	7.890462	2.511554	1.0	6.0	8.0	9.0	21.0

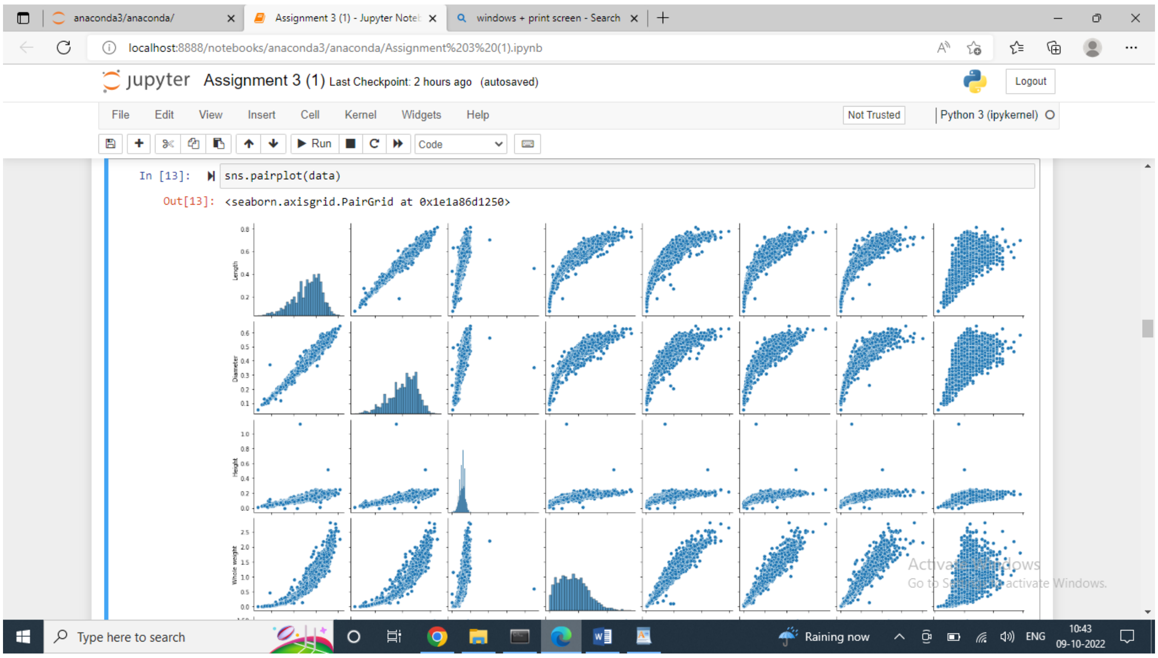
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Note
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter
Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File
Edit
View
Insert
Cell
Kernel
Widgets
Help
Not Trusted
Python 3 (ipykernel)
Run
Code
In [13]:
sns.pairplot(data)
Out[13]:
<seaborn.axisgrid.PairGrid at 0x1e1a86d1250>
Type here to search
Raining now
ENG
10:43
09-10-2022

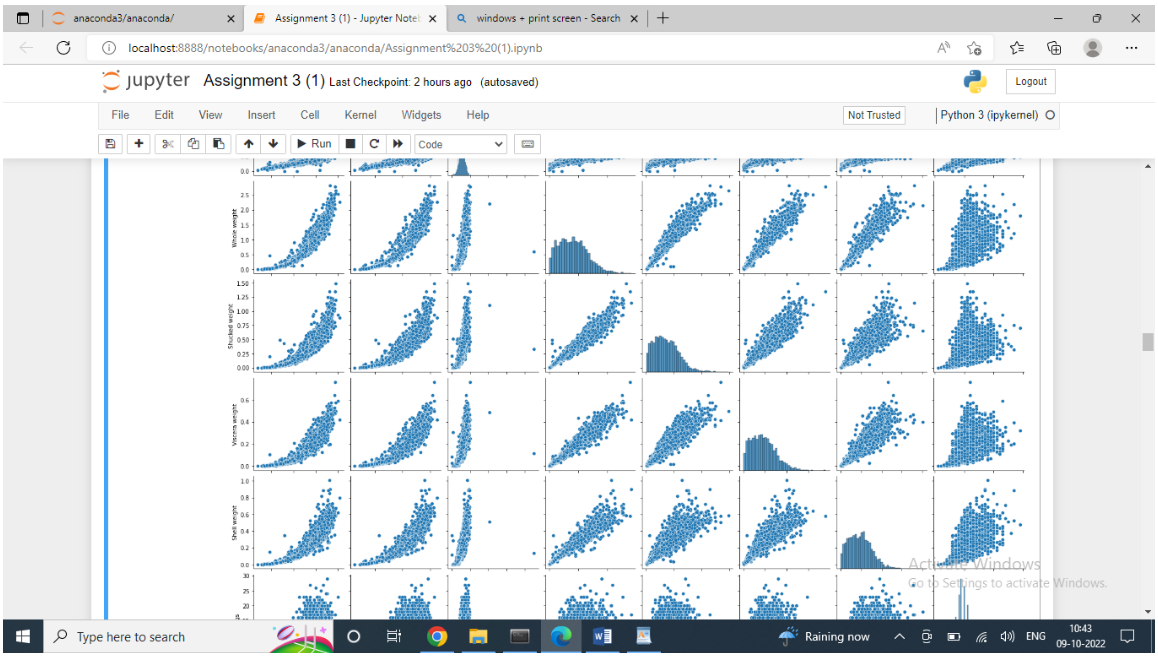
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Noteb
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter
Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File
Edit
View
Insert
Cell
Kernel
Widgets
Help
Not Trusted
Python 3 (ipykernel)
Run
Code
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:43
09-10-2022

jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)

Descriptive statistics

In [14]:
```
data.describe()
```

Out[14]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

Handle The Missing Values

In [15]:
```
data.isnull().any()
```

Out[15]:
```
Sex               False
Length            False
Diameter          False
Height            False
Whole weight      False
Shucked weight    False
```

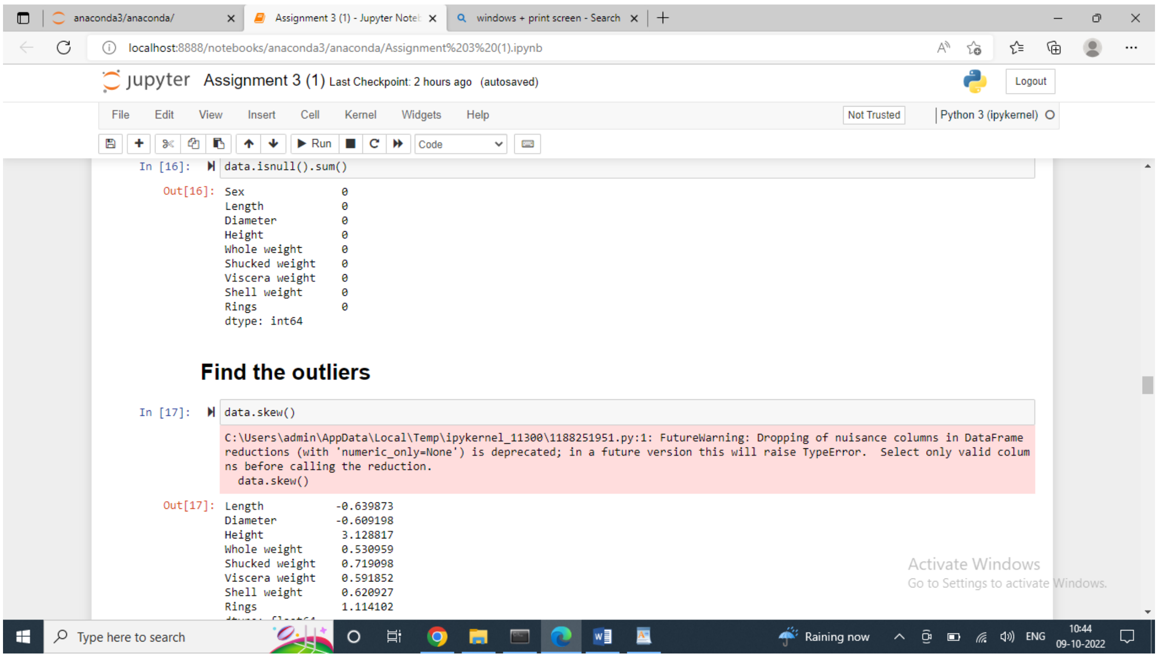

anaconda3/anaconda/
Assignment 3 (1) - Jupyter Notel
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Code
In [16]: data.isnull().sum()
Out[16]: Sex 0
Length 0
Diameter 0
Height 0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings 0
dtype: int64
Find the outliers
In [17]: data.skew()
C:\Users\admin\AppData\Local\Temp\ipykernel_11300\1188251951.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
data.skew()
Out[17]: Length -0.639873
Diameter -0.609198
Height 3.128817
Whole weight 0.530959
Shucked weight 0.719098
Viscera weight 0.591852
Shell weight 0.620927
Rings 1.114102
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:44
09-10-2022

Find the outliers

In [17]:
```
data.skew()
```

```
C:\Users\admin\AppData\Local\Temp\ipykernel_11300\1188251951.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError.  Select only valid columns before calling the reduction.
  data.skew()
```

Out[17]:
```
Length           -0.639873
Diameter         -0.609198
Height            3.128817
Whole weight      0.530959
Shucked weight    0.719098
Viscera weight    0.591852
Shell weight      0.620927
Rings             1.114102
dtype: float64
```

In [18]:
```
sns.boxplot(x=data['Rings'],data=data)
```

Out[18]: <AxesSubplot:xlabel='Rings'>

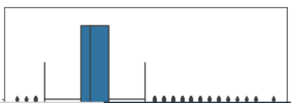

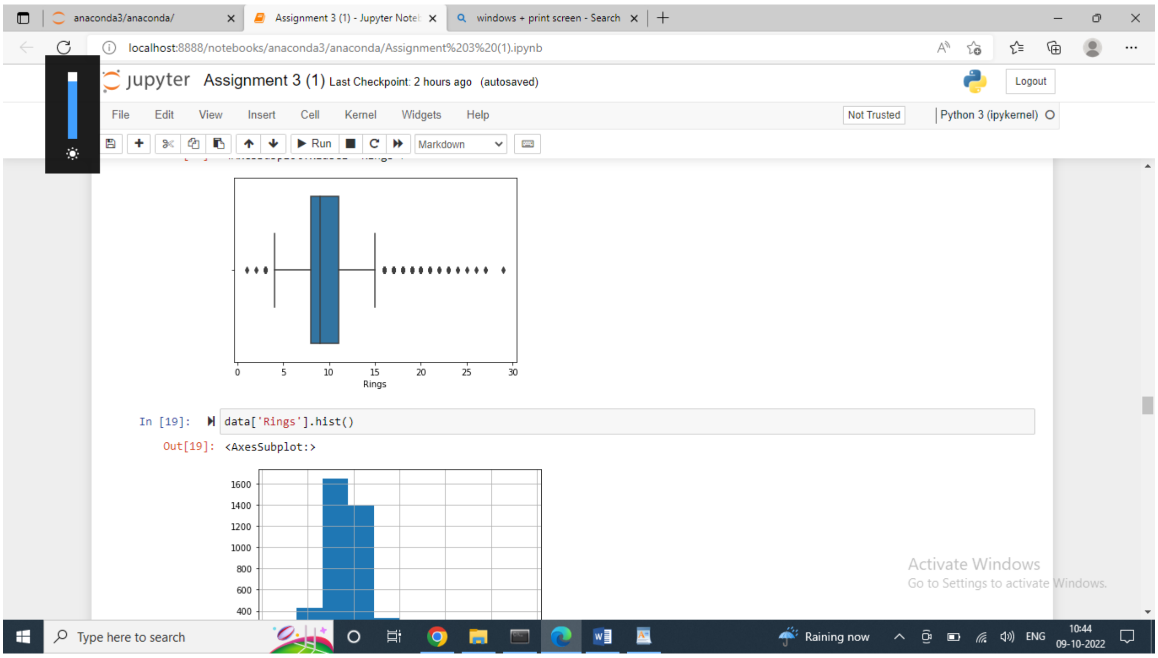

anaconda3/anaconda/
Assignment 3 (1) - Jupyter Notel
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Markdown
0 5 10 15 20 25 30
Rings
In [19]: data['Rings'].hist()
Out[19]: <AxesSubplot:>
1600
1400
1200
1000
800
600
400
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:44
09-10-2022

localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb

jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)

In [20]:
```
print('skewness value of Age: ',data['Rings'].skew())
```

```
skewness value of Age:  1.114101898355677
```

In [21]:
```
Q1 = data['Rings'].quantile(0.25)
Q3 = data['Rings'].quantile(0.75)
IQR = Q3 - Q1
whisker_width = 1.5
lower_whisker = Q1 -(whisker_width*IQR)
upper_whisker = Q3 +(whisker_width*IQR)
data['Rings']=np.where(data['Rings']>upper_whisker,upper_whisker,np.where(data['Rings']<lower_whisker,lower_whisker,data['Rin
```

In [23]:
```
sns.boxplot(x=data['Rings'],data=data)
```

Out[23]: <AxesSubplot:xlabel='Rings'>

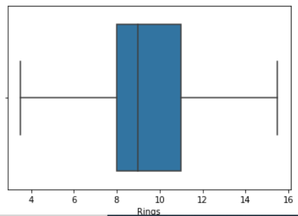

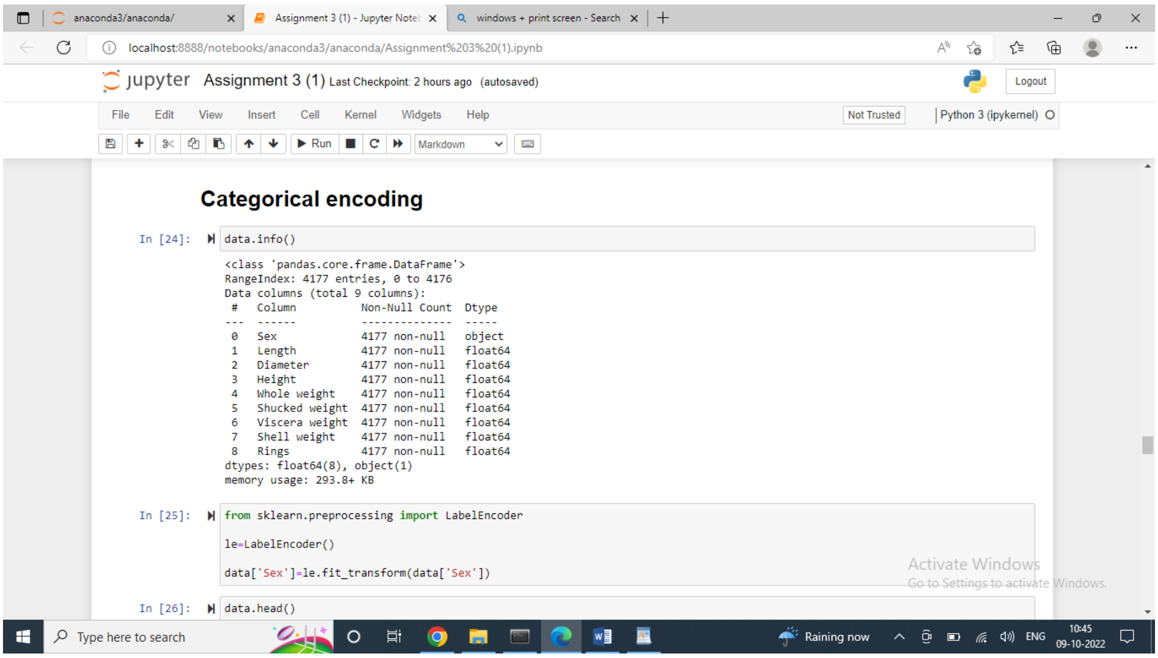

Categorical encoding

In [24]:
```python
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  ------          --------------  -----  
 0   Sex             4177 non-null   object 
 1   Length          4177 non-null   float64
 2   Diameter        4177 non-null   float64
 3   Height          4177 non-null   float64
 4   Whole weight    4177 non-null   float64
 5   Shucked weight  4177 non-null   float64
 6   Viscera weight  4177 non-null   float64
 7   Shell weight    4177 non-null   float64
 8   Rings           4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```

In [25]:
```python
from sklearn.preprocessing import LabelEncoder

le=LabelEncoder()

data['Sex']=le.fit_transform(data['Sex'])
```

In [26]:
```python
data.head()
```

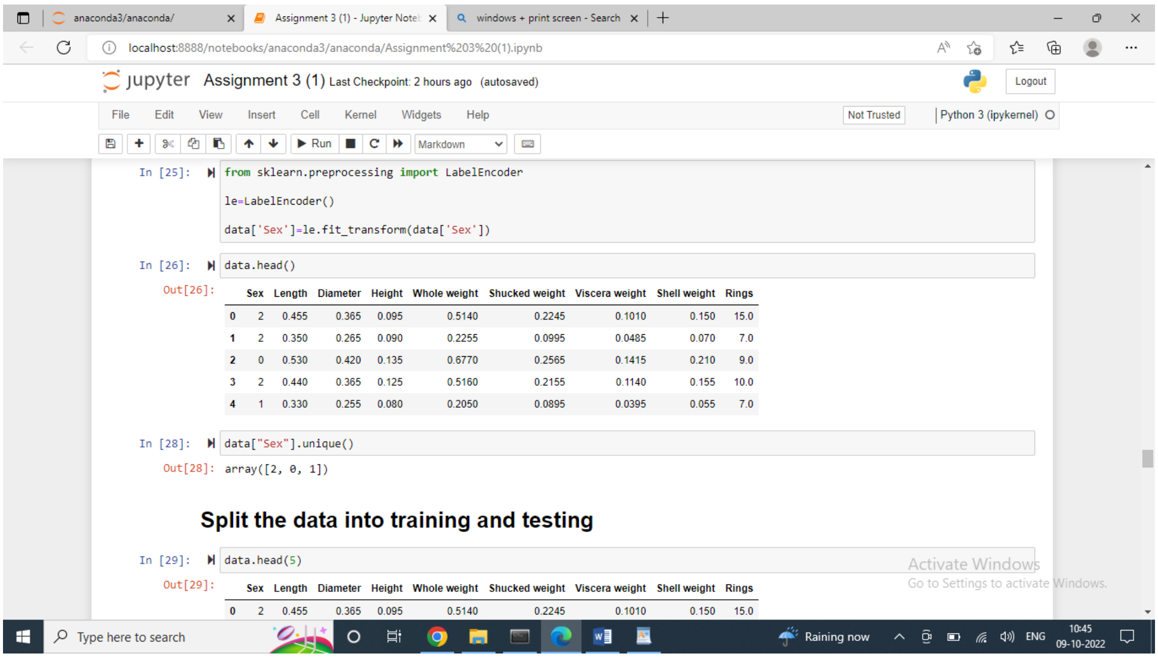

```python
In [25]: from sklearn.preprocessing import LabelEncoder
         le=LabelEncoder()
         data['Sex']=le.fit_transform(data['Sex'])
```

```python
In [26]: data.head()
```

Out[26]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15.0
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.0
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.0
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.0
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.0

```python
In [28]: data["Sex"].unique()
Out[28]: array([2, 0, 1])
```

Split the data into training and testing

```python
In [29]: data.head(5)
```

Out[29]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15.0

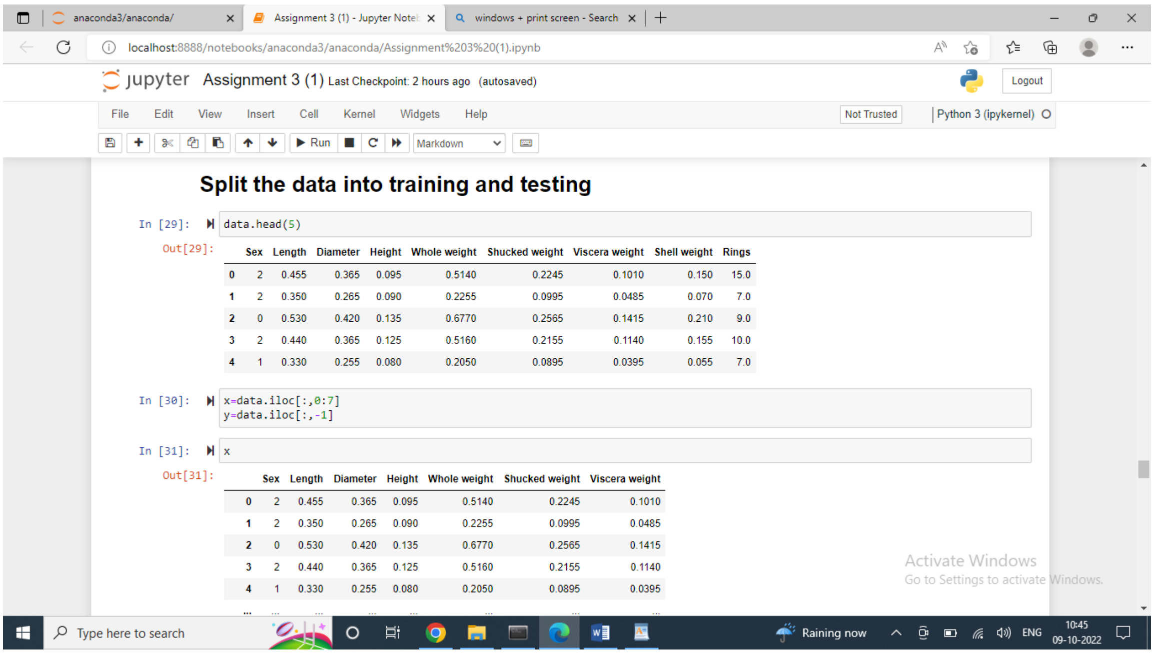

Split the data into training and testing

In [29]:
```
data.head(5)
```

Out[29]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15.0
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.0
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.0
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.0
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.0

In [30]:
```
x=data.iloc[:,0:7]
y=data.iloc[:,-1]
```

In [31]:
```
x
```

Out[31]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395

4177 rows × 7 columns

In [32]:
```
y
```

Out[32]:
```
0       15.0
1        7.0
2        9.0
3       10.0
4        7.0
        ... 
4172    11.0
4173    10.0
4174     9.0
4175    10.0
4176    12.0
Name: Rings, Length: 4177, dtype: float64
```

In [33]:
```
y.shape
```

Out[33]:
```
(4177,)
```

Scale the independent variables

In [36]:
```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
x_scaled = ss.fit_transform(x)
```

Jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)

Out[33]: (4177,)

Scale the independent variables

In [36]:
```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
x_scaled = ss.fit_transform(x)
```

Split the data into training and testing

In [39]:
```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.3, random_state = 1)
```

Build the model

Training the model and testing the model

In [40]:
```
import csv
with open("abalone.csv") as csv_file:
    csv_reader = csv.reader(csv_file)
    data = pd.DataFrame([csv_reader], index = None)
for val in list(data[1]):
    print(val)
['M', '0.455', '0.365', '0.095', '0.514', '0.2245', '0.101', '0.15', '15']
```

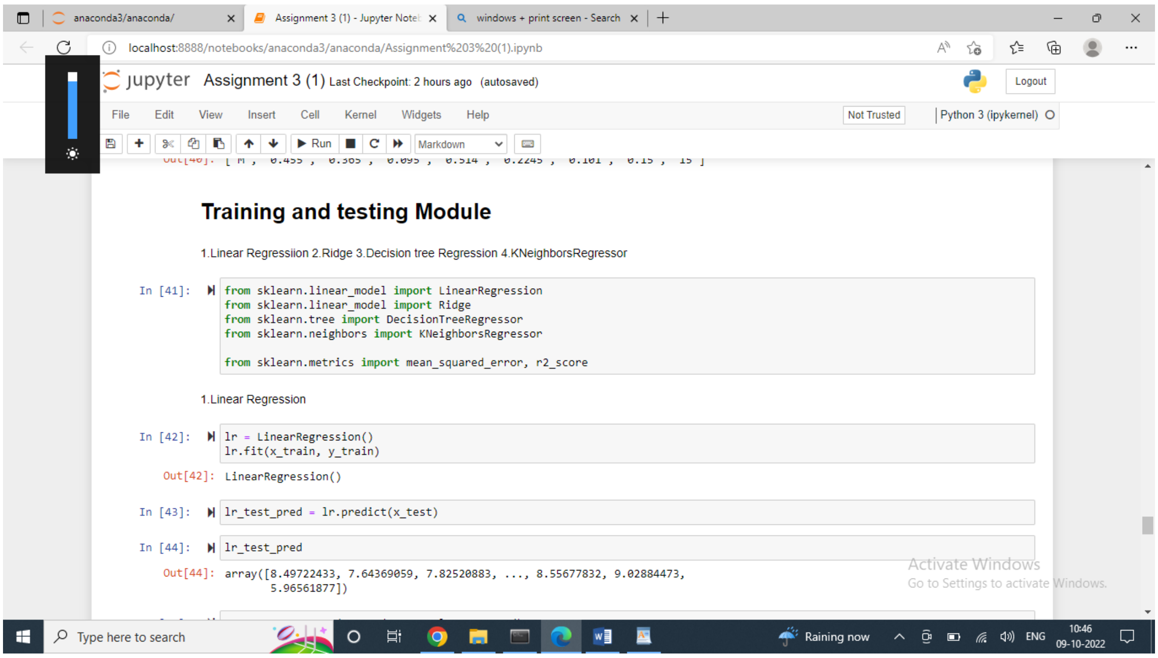
anaconda3/anaconda/
Assignment 3 (1) - Jupyter Note
windows + print screen - Search
localhost:8888/notebooks/anaconda3/anaconda/Assignment%203%20(1).ipynb
jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
Logout
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Run
Markdown
Training and testing Module
1.Linear Regressiion 2.Ridge 3.Decision tree Regression 4.KNeighborsRegressor
In [41]:
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score
1.Linear Regression
In [42]:
lr = LinearRegression()
lr.fit(x_train, y_train)
Out[42]: LinearRegression()
In [43]:
lr_test_pred = lr.predict(x_test)
In [44]:
lr_test_pred
Out[44]: array([8.49722433, 7.64369059, 7.82520883, ..., 8.55677832, 9.02884473,
5.96561877])
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:46
09-10-2022

```
In [42]: lr = LinearRegression()
         lr.fit(x_train, y_train)

Out[42]: LinearRegression()

In [43]: lr_test_pred = lr.predict(x_test)

In [44]: lr_test_pred

Out[44]: array([8.49722433, 7.64369059, 7.82520883, ..., 8.55677832, 9.02884473,
               5.96561877])

In [46]: mse = mean_squared_error(y_test, lr_test_pred)
         print('Mean Squared error of testing Set: %2f'%mse)

         Mean Squared error of testing Set: 3.524602

In [47]: p = r2_score(y_test, lr_test_pred)
         print('R2 Score of testing set:%.2f'%p)

         R2 Score of testing set:0.52
```

1.Ridge

```
In [48]: ridge_mod = Ridge(alpha=0.01, normalize=True)
         ridge_mod.fit(x_train, y_train)
         ridge_mod.fit(x_test, y_test)
```

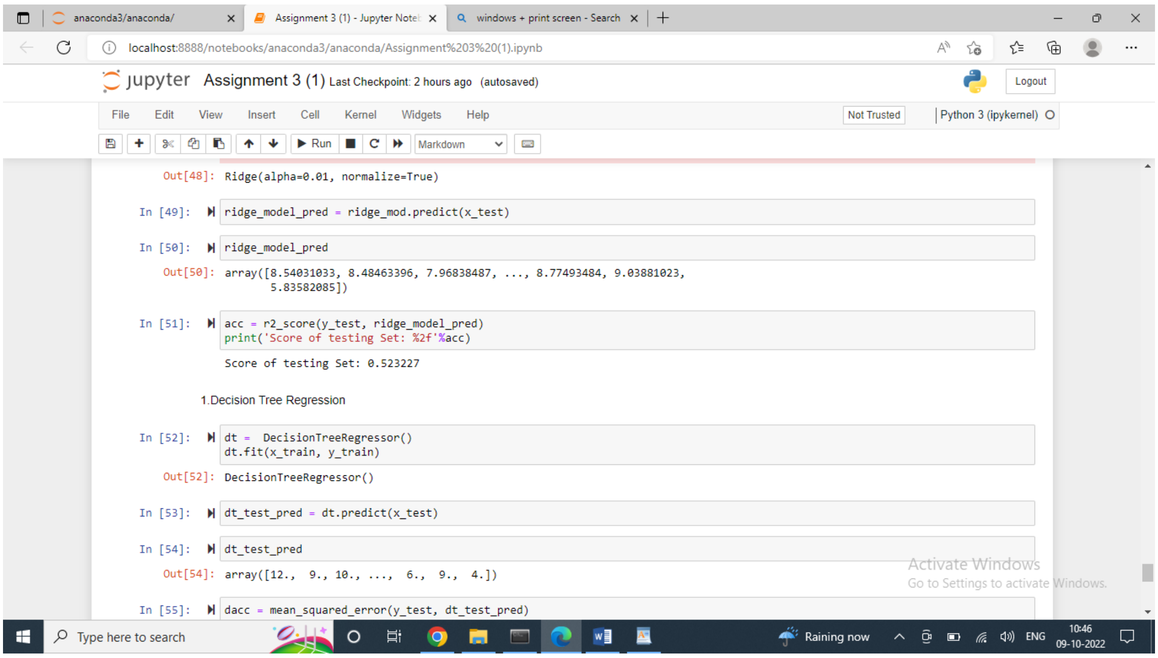

jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)
File Edit View Insert Cell Kernel Widgets Help
Not Trusted
Python 3 (ipykernel)
Out[48]: Ridge(alpha=0.01, normalize=True)
In [49]: ridge_model_pred = ridge_mod.predict(x_test)
In [50]: ridge_model_pred
Out[50]: array([8.54031033, 8.48463396, 7.96838487, ..., 8.77493484, 9.03881023, 5.83582085])
In [51]: acc = r2_score(y_test, ridge_model_pred)
print('Score of testing Set: %2f'%acc)
Score of testing Set: 0.523227
1.Decision Tree Regression
In [52]: dt = DecisionTreeRegressor()
dt.fit(x_train, y_train)
Out[52]: DecisionTreeRegressor()
In [53]: dt_test_pred = dt.predict(x_test)
In [54]: dt_test_pred
Out[54]: array([12., 9., 10., ..., 6., 9., 4.])
In [55]: dacc = mean_squared_error(y_test, dt_test_pred)
Activate Windows
Go to Settings to activate Windows.
Type here to search
Raining now
ENG
10:46
09-10-2022

jupyter Assignment 3 (1) Last Checkpoint: 2 hours ago (autosaved)

In [55]:
```
dacc = mean_squared_error(y_test, dt_test_pred)
print('Mean Squared Error of testing Set: %2f'%dacc)
```

```
Mean Squared Error of testing Set: 6.182217
```

1.KNN Regression

In [56]:
```
knn = KNeighborsRegressor(n_neighbors = 4 )
knn.fit(x_train, y_train)
knn.fit(x_test, y_test)
```

Out[56]:
```
KNeighborsRegressor(n_neighbors=4)
```

In [57]:
```
knn_test_pred = knn.predict(x_test)
```

In [58]:
```
knn_test_pred
```

Out[58]:
```
array([ 8.75,  9.5 , 10.5 , ...,  8.  ,  7.5 ,  5.  ])
```

In [59]:
```
kacc= r2_score(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kacc)
```

```
Score of testing Set: 0.400555
```

In [60]:
```
kmse = mean_squared_error(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kmse)
```

```
Score of testing Set: 2.602460
```